

In [1]:

```

1  ! shred -u setup_google_colab.py
2  ! wget https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/setup_google
3  import setup_google_colab
4  # please, uncomment the week you're working on
5  # setup_google_colab.setup_week1()
6  # setup_google_colab.setup_week2()
7  # setup_google_colab.setup_week2_honor()
8  setup_google_colab.setup_week3()
9  # setup_google_colab.setup_week4()
10 # setup_google_colab.setup_week5()
11 # setup_google_colab.setup_week6()
12 # set tf 1.x for colab
13

```

shred: setup_google_colab.py: failed to open for writing: No such file or directory

--2022-04-11 20:57:10-- https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/setup_google_colab.py (https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/setup_google_colab.py)

Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ...

Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 3636 (3.6K) [text/plain]

Saving to: 'setup_google_colab.py'

setup_google_colab. 100%[=====>] 3.55K --.-KB/s in 0s

2022-04-11 20:57:10 (46.7 MB/s) - 'setup_google_colab.py' saved [3636/3636]

```


*****
102flowers.tgz
*****
imagenet_labels.mat
*****
inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
*****
cifar-10-batches-py.tar.gz
*****
mnist.npz

```

Fine-tuning InceptionV3 for flowers classification

In this task you will fine-tune InceptionV3 architecture for flowers classification task.

InceptionV3 architecture (<https://research.googleblog.com/2016/03/train-your-own-image-classifier-with.html>) (<https://research.googleblog.com/2016/03/train-your-own-image-classifier-with.html>):



Flowers classification dataset (<http://www.robots.ox.ac.uk/~vgg/data/flowers/102/index.html>) consists of 102 flower categories commonly occurring in the United Kingdom. Each class contains between 40 and 258 images:



Import stuff

In [2]:

```
1 import sys
2 sys.path.append("..")
3 import grading
4 import download_utils
```

In [3]:

```
1 # !!! remember to clear session/graph if you rebuild your graph to avoid out-of-
```

In [4]:

```
1 download_utils.link_all_keras_resources()
```

In [5]:

```

1 !pip uninstall keras-nightly
2 !pip uninstall -y tensorflow
3 !pip install keras==2.1.6
4 !pip install tensorflow==1.15.0
5 !pip install install h5py==2.10.0
6
7 import tensorflow as tf
8 import keras
9 from keras import backend as K
10 import numpy as np
11
12 %matplotlib inline
13 import matplotlib.pyplot as plt
14 print(tf.__version__)
15 print(keras.__version__)
16 import cv2 # for image processing
17 from sklearn.model_selection import train_test_split
18 import scipy.io
19 import os
20 import tarfile
21 import keras_utils
22 from keras_utils import reset_tf_session

```

...

Fill in your Coursera token and email

To successfully submit your answers to our grader, please fill in your Coursera submission token and email

In [6]:

```

1 grader = grading.Grader(assignment_key="2v-uxpD7EeeMxQ6FWsz5LA",
2                          all_parts=["wuwWC", "a4FK1", "qRsZ1"])

```

In [7]:

```

1 # token expires every 30 min
2 COURSERA_TOKEN = '4KfJM7vv3brSpvrs'### YOUR TOKEN HERE
3 COURSERA_EMAIL = 'e0321294@u.nus.edu'### YOUR EMAIL HERE

```

Load dataset

Dataset was downloaded for you, it takes 12 min and 400mb. Relevant links (just in case):

- <http://www.robots.ox.ac.uk/~vgg/data/flowers/102/index.html>
(<http://www.robots.ox.ac.uk/~vgg/data/flowers/102/index.html>)
- <http://www.robots.ox.ac.uk/~vgg/data/flowers/102/102flowers.tgz>
(<http://www.robots.ox.ac.uk/~vgg/data/flowers/102/102flowers.tgz>)
- <http://www.robots.ox.ac.uk/~vgg/data/flowers/102/imagelabels.mat>
(<http://www.robots.ox.ac.uk/~vgg/data/flowers/102/imagelabels.mat>)

In [8]:

```
1 # we downloaded them for you, just link them here
2 download_utils.link_week_3_resources()
```

Prepare images for model


In [9]:

```
1 # we will crop and resize input images to IMG_SIZE x IMG_SIZE
2 IMG_SIZE = 250
```

In [10]:

```
1 def decode_image_from_raw_bytes(raw_bytes):
2     img = cv2.imdecode(np.asarray(bytearray(raw_bytes), dtype=np.uint8), 1)
3     img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
4     return img
```

We will take a center crop from each image like this:



In [12]:

```
1 def image_center_crop(img):
2     """
3     Makes a square center crop of an img, which is a [h, w, 3] numpy array.
4     Returns [min(h, w), min(h, w), 3] output with same width and height.
5     For cropping use numpy slicing.
6     """
7     h, w = img.shape[0], img.shape[1]
8     m = min(h, w)
9     cropped_img = img[(h-m)//2:(h+m)//2, (w-m)//2:(w+m)//2, :]
10    ### YOUR CODE HERE
11
12
13    return cropped_img
```

In [13]:

```

1 def prepare_raw_bytes_for_model(raw_bytes, normalize_for_model=True):
2     img = decode_image_from_raw_bytes(raw_bytes) # decode image raw bytes to m
3     img = image_center_crop(img) # take squared center crop
4     img = cv2.resize(img, (IMG_SIZE, IMG_SIZE)) # resize for our model
5     if normalize_for_model:
6         img = img.astype("float32") # prepare for normalization
7         img = keras.applications.inception_v3.preprocess_input(img) # normalize
8     return img

```

In [14]:

```

1 # reads bytes directly from tar by filename (slow, but ok for testing, takes ~6
2 def read_raw_from_tar(tar_fn, fn):
3     with tarfile.open(tar_fn) as f:
4         m = f.getmember(fn)
5         return f.extractfile(m).read()

```

In [15]:

```

1 # test cropping
2 raw_bytes = read_raw_from_tar("102flowers.tgz", "jpg/image_00001.jpg")
3
4 img = decode_image_from_raw_bytes(raw_bytes)
5 print(img.shape)
6 plt.imshow(img)
7 plt.show()
8
9 img = prepare_raw_bytes_for_model(raw_bytes, normalize_for_model=False)
10 print(img.shape)
11 plt.imshow(img)
12 plt.show()

```

...

In [16]:

```

1 ## GRADED PART, DO NOT CHANGE!
2 # Test image preparation for model
3 prepared_img = prepare_raw_bytes_for_model(read_raw_from_tar("102flowers.tgz", '
4 grader.set_answer("qRsZ1", list(prepared_img.shape) + [np.mean(prepared_img), np

```

In [17]:

```

1 # you can make submission with answers so far to check yourself at this stage
2 grader.submit(COURSE_EMAIL, COURSE_TOKEN)

```

...

Prepare for training

In [18]:

```
1 # read all filenames and labels for them
2
3 # read filenames firectly from tar
4 def get_all_filenames(tar_fn):
5     with tarfile.open(tar_fn) as f:
6         return [m.name for m in f.getmembers() if m.isfile()]
7
8 all_files = sorted(get_all_filenames("102flowers.tgz")) # list all files in tar
9 all_labels = scipy.io.loadmat('imagelabels.mat')['labels'][0] - 1 # read class
10 # all_files and all_labels are aligned now
11 N_CLASSES = len(np.unique(all_labels))
12 print(N_CLASSES)
```

...

In [19]:

```
1 # split into train/test
2 tr_files, te_files, tr_labels, te_labels = \
3     train_test_split(all_files, all_labels, test_size=0.2, random_state=42, stratify=all_labels)
```

In [20]:

```
1 # will yield raw image bytes from tar with corresponding label
2 def raw_generator_with_label_from_tar(tar_fn, files, labels):
3     label_by_fn = dict(zip(files, labels))
4     with tarfile.open(tar_fn) as f:
5         while True:
6             m = f.next()
7             if m is None:
8                 break
9             if m.name in label_by_fn:
10                 yield f.extractfile(m).read(), label_by_fn[m.name]
```

In [21]:

```

1  # batch generator
2  BATCH_SIZE = 32
3
4  def batch_generator(items, batch_size):
5      """
6      Implement batch generator that yields items in batches of size batch_size.
7      There's no need to shuffle input items, just chop them into batches.
8      Remember about the last batch that can be smaller than batch_size!
9      Input: any iterable (list, generator, ...). You should do `for item in items`
10         In case of generator you can pass through your items only once!
11      Output: In output yield each batch as a list of items.
12      """
13      #
14      batch = []
15      for i, item in enumerate(items):
16          batch.append(item)
17          if i % batch_size == batch_size-1:
18              yield batch
19              batch = []
20      if batch[0]:
21          yield [item for item in batch if item]
22      ### YOUR CODE HERE

```

In [22]:

```

1  ## GRADED PART, DO NOT CHANGE!
2  # Test batch generator
3  def _test_items_generator():
4      for i in range(10):
5          yield i
6
7  grader.set_answer("a4FK1", list(map(lambda x: len(x), batch_generator(_test_items_generator()))))

```

In [23]:

```

1  # you can make submission with answers so far to check yourself at this stage
2  grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)

```

...

In [24]:

```

1  def train_generator(files, labels):
2      while True: # so that Keras can loop through this as long as it wants
3          for batch in batch_generator(raw_generator_with_label_from_tar(
4              "102flowers.tgz", files, labels), BATCH_SIZE):
5              # prepare batch images
6              batch_imgs = []
7              batch_targets = []
8              for raw, label in batch:
9                  img = prepare_raw_bytes_for_model(raw)
10                 batch_imgs.append(img)
11                 batch_targets.append(label)
12             # stack images into 4D tensor [batch_size, img_size, img_size, 3]
13             batch_imgs = np.stack(batch_imgs, axis=0)
14             # convert targets into 2D tensor [batch_size, num_classes]
15             batch_targets = keras.utils.np_utils.to_categorical(batch_targets, N)
16             yield batch_imgs, batch_targets

```

In [25]:

```

1 # test training generator
2 for _ in train_generator(tr_files, tr_labels):
3     print(_[0].shape, _[1].shape)
4     plt.imshow(np.clip(_[0][0] / 2. + 0.5, 0, 1))
5     break

```

...

Training

You cannot train such a huge architecture from scratch with such a small dataset.

But using fine-tuning of last layers of pre-trained network you can get a pretty good classifier very quickly.

In [26]:

```

1 # remember to clear session if you start building graph from scratch!
2 s = reset_tf_session()
3 # don't call K.set_learning_phase() !!! (otherwise will enable dropout in train,

```

...

In [27]:

```

1 def inception(use_imagenet=True):
2     # load pre-trained model graph, don't add final layer
3     model = keras.applications.InceptionV3(include_top=False, input_shape=(IMG_S
4                                         weights='imagenet' if use_imagenet else
5     # add global pooling just like in InceptionV3
6     new_output = keras.layers.GlobalAveragePooling2D()(model.output)
7     # add new dense layer for our labels
8     new_output = keras.layers.Dense(N_CLASSES, activation='softmax')(new_output)
9     model = keras.engine.training.Model(model.inputs, new_output)
10    return model

```

In [28]:

```
1 model = inception()
```

...

In [29]:

```
1 model.summary()
```

...

In [30]:

```

1 # how many layers our model has
2 print(len(model.layers))

```

...

In [31]:

```

1 # set all layers trainable by default
2 for layer in model.layers:
3     layer.trainable = True
4     if isinstance(layer, keras.layers.BatchNormalization):
5         # we do aggressive exponential smoothing of batch norm
6         # parameters to faster adjust to our new dataset
7         layer.momentum = 0.9
8
9 # fix deep layers (fine-tuning only last 50)
10 for layer in model.layers[:-50]:
11     # fix all but batch norm layers, because we needed to update moving averages
12     if not isinstance(layer, keras.layers.BatchNormalization):
13         layer.trainable = False

```

In [32]:

```

1 # compile new model
2 model.compile(
3     loss='categorical_crossentropy', # we train 102-way classification
4     optimizer=keras.optimizers.adamax(lr=1e-2), # we can take big lr here because
5     metrics=['accuracy'] # report accuracy during training
6 )

```

...

In [33]:

```

1 # we will save model checkpoints to continue training in case of kernel death
2 model_filename = 'flowers.{0:03d}.hdf5'
3 last_finished_epoch = None
4
5 ##### uncomment below to continue training from model checkpoint
6 ##### fill `last_finished_epoch` with your latest finished epoch
7 # from keras.models import load_model
8 # s = reset_tf_session()
9 # last_finished_epoch = 10
10 # model = load_model(model_filename.format(last_finished_epoch))

```

Training takes **2 hours**. You're aiming for ~0.93 validation accuracy.

In [34]:

```

1 # fine tune for 2 epochs (full passes through all training data)
2 # we make 2*8 epochs, where epoch is 1/8 of our training data to see progress more
3 model.fit_generator(
4     train_generator(tr_files, tr_labels),
5     steps_per_epoch=len(tr_files) // BATCH_SIZE // 8,
6     epochs=2 * 8,
7     validation_data=train_generator(te_files, te_labels),
8     validation_steps=len(te_files) // BATCH_SIZE // 4,
9     callbacks=[keras_utils.TqdmProgressCallback(),
10                keras_utils.ModelSaveCallback(model_filename)],
11     verbose=0,
12     initial_epoch=last_finished_epoch or 0
13 )

```

...

In [35]:

```
1  ## GRADED PART, DO NOT CHANGE!
2  # Accuracy on validation set
3  test_accuracy = model.evaluate_generator(
4      train_generator(te_files, te_labels),
5      len(te_files) // BATCH_SIZE // 2
6  )[1]
7  grader.set_answer("wuwWC", test_accuracy)
8  print(test_accuracy)
```

...

In [36]:

```
1  # you can make submission with answers so far to check yourself at this stage
2  grader.submit(COURSE_EMAIL, COURSE_TOKEN)
```

...

That's it! Congratulations!

What you've done:

- prepared images for the model
- implemented your own batch generator
- fine-tuned the pre-trained model