

```

! shred -u setup_google_colab.py
! wget https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/setup_google_co
import setup_google_colab
# please, uncomment the week you're working on
# setup_google_colab.setup_week1()
# setup_google_colab.setup_week2()
# setup_google_colab.setup_week2_honor()
# setup_google_colab.setup_week3()
setup_google_colab.setup_week4()
# setup_google_colab.setup_week5()
# setup_google_colab.setup_week6()
# set tf 1.x for colab

# set tf 1.x for colab
%tensorflow_version 1.x

```

```

--2022-04-18 06:55:26-- https://raw.githubusercontent.com/hse-aml/intro-to-dl
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.109
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.109
HTTP request sent, awaiting response... 200 OK
Length: 3636 (3.6K) [text/plain]
Saving to: 'setup_google_colab.py'

```

```

setup_google_colab. 100%[=====>] 3.55K --.-KB/s in 0s

```

```

2022-04-18 06:55:26 (40.6 MB/s) - 'setup_google_colab.py' saved [3636/3636]

```

```

*****
lfw-deepfunneled.tgz
*****
lfw.tgz
*****
lfw_attributes.txt
TensorFlow 1.x selected.

```

Denoising Autoencoders And Where To Find Them

Today we're going to train deep autoencoders and apply them to faces and similar images search.

Our new test subjects are human faces from the [lfw dataset](#).

▼ Import stuff

```

import sys
sys.path.append("..")
import grading

```

```

!pip uninstall keras-nightly
!pip uninstall -y tensorflow
!pip install tensorflow==1.15.0
!pip install keras==2.1.6
!pip install install h5py==2.10.0

import tensorflow as tf
import keras, keras.layers as L, keras.backend as K
import numpy as np
from sklearn.model_selection import train_test_split
from lfw_dataset import load_lfw_dataset
%matplotlib inline
import matplotlib.pyplot as plt
import download_utils
import keras_utils
import numpy as np
from keras_utils import reset_tf_session

```

WARNING: Skipping keras-nightly as it is not installed.

Found existing installation: tensorflow 1.15.0

Uninstalling tensorflow-1.15.0:

Successfully uninstalled tensorflow-1.15.0

Collecting tensorflow==1.15.0

Using cached tensorflow-1.15.0-cp37-cp37m-manylinux2010_x86_64.whl (412.3 MB)

Requirement already satisfied: tensorflow-estimator==1.15.1 in /tensorflow-1.15.0 (412.3 MB)

Requirement already satisfied: tensorboard<1.16.0,>=1.15.0 in /tensorflow-1.15.0 (412.3 MB)

Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.7/dist-packages (3.6.1)

Requirement already satisfied: keras-applications>=1.0.8 in /tensorflow-1.15.0 (412.3 MB)

Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.7/dist-packages (0.6.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.7/dist-packages (1.1.0)

Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.7/dist-packages (1.8.6)

Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.7/dist-packages (0.1.6)

Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.7/dist-packages (1.0.5)

Requirement already satisfied: gast==0.2.2 in /usr/local/lib/python3.7/dist-packages (0.2.2)

Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.7/dist-packages (1.11.1)

Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.7/dist-packages (0.26)

Requirement already satisfied: numpy<2.0,>=1.16.0 in /usr/local/lib/python3.7/dist-packages (1.16.0)

Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.7/dist-packages (1.10.0)

Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.7/dist-packages (0.7.0)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.7/dist-packages (2.3.2)

Requirement already satisfied: h5py in /usr/local/lib/python3.7/dist-packages (2.10.0)

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (2.6.8)

Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packages (41.0.0)

Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packages (0.11.15)

Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (4.4)

Requirement already satisfied: typing-extensions>=3.6.4 in /usr/local/lib/python3.7/dist-packages (3.6.4)

Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (0.5)

Installing collected packages: tensorflow

ERROR: pip's dependency resolver does not currently take into account all the dependencies of tensorflow==1.15.0, but you have tensorflow 1.15.0 which is already installed.

Successfully installed tensorflow-1.15.0

Collecting keras==2.1.6

Using cached Keras-2.1.6-py2.py3-none-any.whl (339 kB)

Requirement already satisfied: h5py in /usr/local/lib/python3.7/dist-packages (2.10.0)

Requirement already satisfied: scipy>=0.14 in /usr/local/lib/python3.7/dist-packages (0.14)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (5.4.1)

Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.7/dist-packages (1.9.0)

```
Requirement already satisfied: numpy>=1.9.1 in /usr/local/lib/python3.7/dist-packages
Installing collected packages: keras
  Attempting uninstall: keras
    Found existing installation: Keras 2.0.6
    Uninstalling Keras-2.0.6:
      Successfully uninstalled Keras-2.0.6
Successfully installed keras-2.1.6
Requirement already satisfied: install in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: h5py==2.10.0 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: numpy>=1.7 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (
Using TensorFlow backend.
```

```
# !!! remember to clear session/graph if you rebuild your graph to avoid out-of-mem
```

▼ Load dataset

Dataset was downloaded for you. Relevant links (just in case):

- http://www.cs.columbia.edu/CAVE/databases/pubfig/download/lfw_attributes.txt
- <http://vis-www.cs.umass.edu/lfw/lfw-deepfunneled.tgz>
- <http://vis-www.cs.umass.edu/lfw/lfw.tgz>

```
# we downloaded them for you, just link them here
download_utils.link_week_4_resources()
```

```
# load images
X, attr = load_lfw_dataset(use_raw=True, dimx=32, dimy=32)
IMG_SHAPE = X.shape[1:]

# center images
X = X.astype('float32') / 255.0 - 0.5

# split
X_train, X_test = train_test_split(X, test_size=0.1, random_state=42)
```

```
*****
```

```
def show_image(x):
    plt.imshow(np.clip(x + 0.5, 0, 1))
```

```
plt.title('sample images')
```

```
for i in range(6):
    plt.subplot(2,3,i+1)
    show_image(X[i])
```

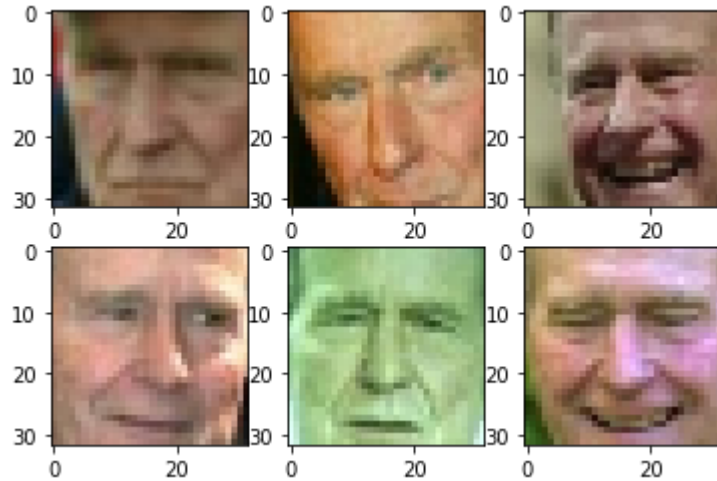
```
print("X shape:", X.shape)
print("attr shape:", attr.shape)
```

```
# try to free memory  
del X  
import gc  
gc.collect()
```

X shape: (13143, 32, 32, 3)

attr shape: (13143, 73)


1385



Autoencoder architecture

Let's design autoencoder as two sequential keras models: the encoder and decoder respectively.

We will then use symbolic API to apply and train these models.



▼ First step: PCA


Principal Component Analysis is a popular dimensionality reduction method.

Under the hood, PCA attempts to decompose object-feature matrix X into two smaller matrices: W and \hat{W} minimizing *mean squared error*:

$$\|(XW)\hat{W} - X\|_2^2 \rightarrow_{W, \hat{W}} \min$$

- $X \in \mathbb{R}^{n \times m}$ - object matrix (**centered**);
- $W \in \mathbb{R}^{m \times d}$ - matrix of direct transformation;
- $\hat{W} \in \mathbb{R}^{d \times m}$ - matrix of reverse transformation;
- n samples, m original dimensions and d target dimensions;

In geometric terms, we want to find d axes along which most of variance occurs. The "natural" axes, if you wish.



PCA can also be seen as a special case of an autoencoder.

- **Encoder:** $X \rightarrow \text{Dense}(d \text{ units}) \rightarrow \text{code}$
- **Decoder:** $\text{code} \rightarrow \text{Dense}(m \text{ units}) \rightarrow X$

Where Dense is a fully-connected layer with linear activation: $f(X) = W \cdot X + \vec{b}$

Note: the bias term in those layers is responsible for "centering" the matrix i.e. subtracting mean.

```
def build_pca_autoencoder(img_shape, code_size):
    """
    Here we define a simple linear autoencoder as described above.
    We also flatten and un-flatten data to be compatible with image shapes
    """

    encoder = keras.models.Sequential()
    encoder.add(L.InputLayer(img_shape))
    encoder.add(L.Flatten())                #flatten image to vector
    encoder.add(L.Dense(code_size))         #actual encoder

    decoder = keras.models.Sequential()
    decoder.add(L.InputLayer((code_size,)))
    decoder.add(L.Dense(np.prod(img_shape))) #actual decoder, height*width*3 units
    decoder.add(L.Reshape(img_shape))       #un-flatten

    return encoder, decoder
```

Meld them together into one model:

```
s = reset_tf_session()

encoder, decoder = build_pca_autoencoder(IMG_SHAPE, code_size=32)

inp = L.Input(IMG_SHAPE)
code = encoder(inp)
reconstruction = decoder(code)

autoencoder = keras.models.Model(inputs=inp, outputs=reconstruction)
autoencoder.compile(optimizer='adamax', loss='mse')

autoencoder.fit(x=X_train, y=X_train, epochs=15,
```

```
validation_data=[X_test, X_test],
verbose=True)#remove the callbacks
```

Train on 11828 samples, validate on 1315 samples

Epoch 1/15

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/t

11828/11828 [=====] - 3s 275us/step - loss: 0.0123 -

Epoch 2/15

11828/11828 [=====] - 3s 239us/step - loss: 0.0077 -

Epoch 3/15

11828/11828 [=====] - 2s 176us/step - loss: 0.0069 -

Epoch 4/15

11828/11828 [=====] - 2s 191us/step - loss: 0.0067 -

Epoch 5/15

11828/11828 [=====] - 2s 184us/step - loss: 0.0067 -

Epoch 6/15

11828/11828 [=====] - 2s 178us/step - loss: 0.0067 -

Epoch 7/15

11828/11828 [=====] - 2s 175us/step - loss: 0.0067 -

Epoch 8/15

11828/11828 [=====] - 2s 182us/step - loss: 0.0067 -

Epoch 9/15

11828/11828 [=====] - 2s 174us/step - loss: 0.0067 -

Epoch 10/15

11828/11828 [=====] - 2s 174us/step - loss: 0.0067 -

Epoch 11/15

11828/11828 [=====] - 2s 176us/step - loss: 0.0067 -

Epoch 12/15

11828/11828 [=====] - 2s 184us/step - loss: 0.0067 -

Epoch 13/15

11828/11828 [=====] - 2s 172us/step - loss: 0.0067 -

Epoch 14/15

11828/11828 [=====] - 2s 174us/step - loss: 0.0067 -

Epoch 15/15

11828/11828 [=====] - 2s 181us/step - loss: 0.0067 -

<keras.callbacks.History at 0x7f5c23125690>

```
def visualize(img,encoder,decoder):
```

```
    """Draws original, encoded and decoded images"""
```

```
    code = encoder.predict(img[None])[0] # img[None] is the same as img[np.newaxis]
```

```
    reco = decoder.predict(code[None])[0]
```

```
    plt.subplot(1,3,1)
```

```
    plt.title("Original")
```

```
    show_image(img)
```

```
    plt.subplot(1,3,2)
```

```
    plt.title("Code")
```

```
    plt.imshow(code.reshape([code.shape[-1]//2,-1]))
```

```
    plt.subplot(1,3,3)
```

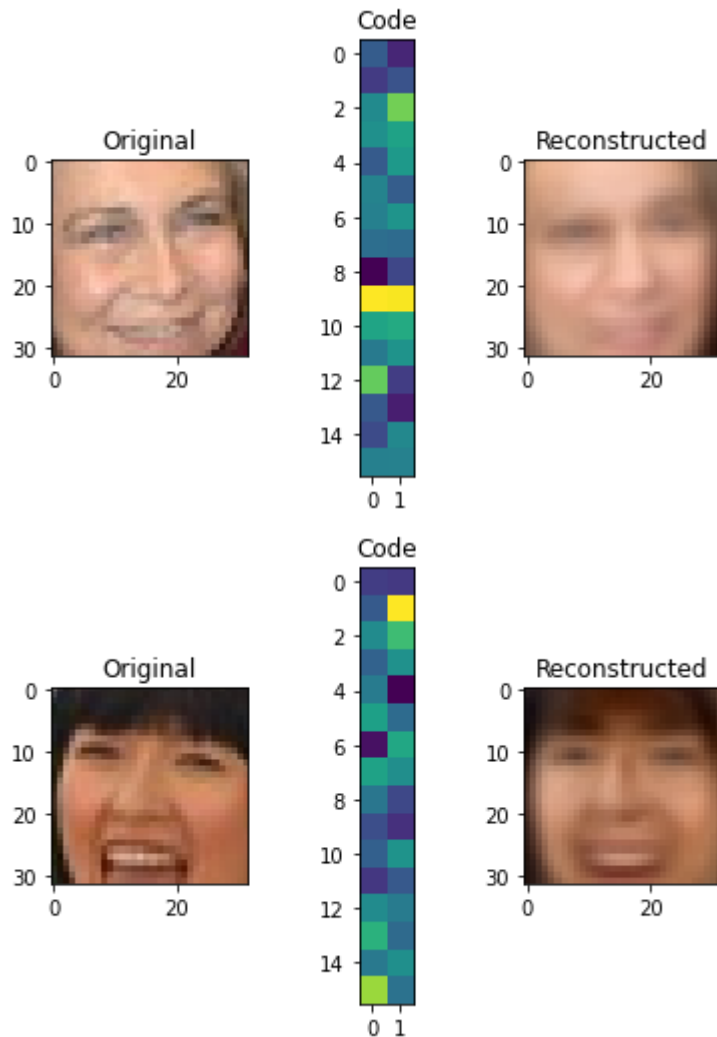
```
    plt.title("Reconstructed")
```

```
show_image(reco)
plt.show()
```

```
score = autoencoder.evaluate(X_test,X_test,verbose=0)
print("PCA MSE:", score)

for i in range(5):
    img = X_test[i]
    visualize(img,encoder,decoder)
```


PCA MSE: 0.006608356407953306



▼ Going deeper: convolutional autoencoder

PCA is neat but surely we can do better. This time we want you to build a deep convolutional autoencoder by... stacking more layers.

Encoder

The **encoder** part is pretty standard, we stack convolutional and pooling layers and finish with a dense layer to get the representation of desirable size (`code_size`).

We recommend to use `activation='elu'` for all convolutional and dense layers.

We recommend to repeat (conv, pool) 4 times with kernel size (3, 3), `padding='same'` and the following numbers of output channels: 32, 64, 128, 256.

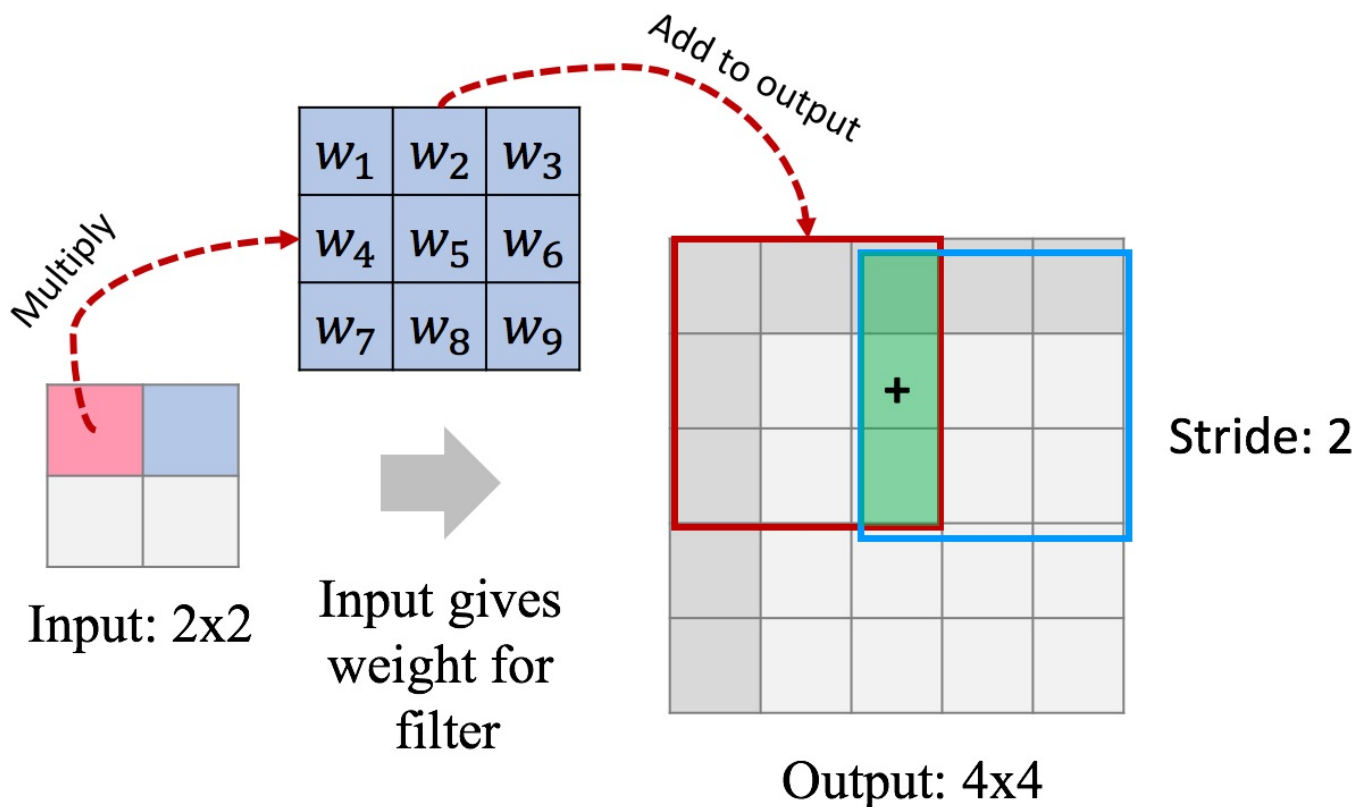
Remember to flatten (`L.Flatten()`) output before adding the last dense layer!

Decoder

For **decoder** we will use so-called "transpose convolution".

Traditional convolutional layer takes a patch of an image and produces a number (patch \rightarrow number). In "transpose convolution" we want to take a number and produce a patch of an image (number \rightarrow patch). We need this layer to "undo" convolutions in encoder. We had a glimpse of it during week 3 (watch [this video](#) starting at 5:41).

Here's how "transpose convolution" works:



In this example we use a stride of 2 to produce 4x4 output, this way we "undo" pooling as well. Another way to think about it: we "undo" convolution with stride 2 (which is similar to conv + pool).

You can add "transpose convolution" layer in Keras like this:

```
L.Conv2DTranspose(filters=?, kernel_size=(3, 3), strides=2, activation='elu', padding='sam
```

Our decoder starts with a dense layer to "undo" the last layer of encoder. Remember to reshape its output to "undo" `L.Flatten()` in encoder.

Now we're ready to undo (conv, pool) pairs. For this we need to stack 4 `L.Conv2DTranspose` layers with the following numbers of output channels: 128, 64, 32, 3. Each of these layers will learn to "undo" (conv, pool) pair in encoder. For the last `L.Conv2DTranspose` layer use `activation=None` because that is our final image.

```
# Let's play around with transpose convolution on examples first
def test_conv2d_transpose(img_size, filter_size):
    print("Transpose convolution test for img_size={}, filter_size={}".format(img_

    x = (np.arange(img_size ** 2, dtype=np.float32) + 1).reshape((1, img_size, img_
    f = (np.ones(filter_size ** 2, dtype=np.float32)).reshape((filter_size, filter_

    s = reset_tf_session()

    conv = tf.nn.conv2d_transpose(x, f,
                                   output_shape=(1, img_size * 2, img_size * 2, 1),
                                   strides=[1, 2, 2, 1],
                                   padding='SAME')

    result = s.run(conv)
    print("input:")
    print(x[0, :, :, 0])
    print("filter:")
    print(f[:, :, 0, 0])
    print("output:")
    print(result[0, :, :, 0])
    s.close()

test_conv2d_transpose(img_size=2, filter_size=2)
test_conv2d_transpose(img_size=2, filter_size=3)
test_conv2d_transpose(img_size=4, filter_size=2)
test_conv2d_transpose(img_size=4, filter_size=3)
```

Transpose convolution test for img_size=2, filter_size=2:

input:

```
[[1. 2.]
 [3. 4.]]
```

filter:

```
[[1. 1.]
 [1. 1.]]
```

output:

```
[[1. 1. 2. 2.]
 [1. 1. 2. 2.]
 [3. 3. 4. 4.]
 [3. 3. 4. 4.]]
```

Transpose convolution test for img_size=2, filter_size=3:

input:

```
[[1. 2.]
 [3. 4.]]
```

filter:

```
[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
```

output:

```

[[ 1.  1.  3.  2.]
 [ 1.  1.  3.  2.]
 [ 4.  4. 10.  6.]
 [ 3.  3.  7.  4.]]
Transpose convolution test for img_size=4, filter_size=2:
input:
[[ 1.  2.  3.  4.]
 [ 5.  6.  7.  8.]
 [ 9. 10. 11. 12.]
 [13. 14. 15. 16.]]
filter:
[[1. 1.]
 [1. 1.]]
output:
[[ 1.  1.  2.  2.  3.  3.  4.  4.]
 [ 1.  1.  2.  2.  3.  3.  4.  4.]
 [ 5.  5.  6.  6.  7.  7.  8.  8.]
 [ 5.  5.  6.  6.  7.  7.  8.  8.]
 [ 9.  9. 10. 10. 11. 11. 12. 12.]
 [ 9.  9. 10. 10. 11. 11. 12. 12.]
 [13. 13. 14. 14. 15. 15. 16. 16.]
 [13. 13. 14. 14. 15. 15. 16. 16.]]
Transpose convolution test for img_size=4, filter_size=3:
input:
[[ 1.  2.  3.  4.]
 [ 5.  6.  7.  8.]
 [ 9. 10. 11. 12.]
 [13. 14. 15. 16.]]
filter:
[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
output:
[[ 1.  1.  3.  2.  5.  3.  7.  4.]
 [ 1.  1.  3.  2.  5.  3.  7.  4.]
 [ 6.  6. 14.  8. 18. 10. 22. 12.]

```

```

def build_deep_autoencoder(img_shape, code_size):
    """PCA's deeper brother. See instructions above. Use `code_size` in layer defin
    H,W,C = img_shape

    # encoder
    encoder = keras.models.Sequential()
    encoder.add(L.InputLayer(img_shape))

    ### YOUR CODE HERE: define encoder as per instructions above ###
    encoder.add(L.Conv2D(filters = 32, kernel_size = (3, 3), padding = 'same', acti
    encoder.add(L.MaxPooling2D())
    encoder.add(L.Conv2D(filters = 64, kernel_size = (3, 3), padding = 'same', acti
    encoder.add(L.MaxPooling2D())
    encoder.add(L.Conv2D(filters = 128, kernel_size = (3, 3), padding = 'same', act
    encoder.add(L.MaxPooling2D())
    encoder.add(L.Conv2D(filters = 256, kernel_size = (3, 3), padding = 'same', act
    encoder.add(L.MaxPooling2D())
    encoder.add(L.Flatten())
    encoder.add(L.Dense(code_size))

```

```
# decoder
decoder = keras.models.Sequential()
decoder.add(L.InputLayer((code_size,)))

### YOUR CODE HERE: define decoder as per instructions above ###
decoder.add(L.Dense(2 * 2 * 256))
decoder.add(L.Reshape((2,2,256)))
decoder.add(L.Conv2DTranspose(filters = 128, kernel_size = (3, 3), strides = 2,
decoder.add(L.Conv2DTranspose(filters = 64, kernel_size = (3, 3), strides = 2,
decoder.add(L.Conv2DTranspose(filters = 32, kernel_size = (3, 3), strides = 2,
decoder.add(L.Conv2DTranspose(filters = 3, kernel_size = (3, 3), strides = 2,

return encoder, decoder
```

```
# Check autoencoder shapes along different code_sizes
get_dim = lambda layer: np.prod(layer.output_shape[1:])
for code_size in [1,8,32,128,512]:
    s = reset_tf_session()
    encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=code_size)
    print("Testing code size %i" % code_size)
    assert encoder.output_shape[1]==(code_size,),"encoder must output a code of re
    assert decoder.output_shape[1]==IMG_SHAPE, "decoder must output an image of
    assert len(encoder.trainable_weights)>=6, "encoder must contain at least 3
    assert len(decoder.trainable_weights)>=6, "decoder must contain at least 3

    for layer in encoder.layers + decoder.layers:
        assert get_dim(layer) >= code_size, "Encoder layer %s is smaller than bottl

print("All tests passed!")
s = reset_tf_session()
```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/t

```
Testing code size 1
Testing code size 8
Testing code size 32
Testing code size 128
Testing code size 512
All tests passed!
```

```
# Look at encoder and decoder shapes.
# Total number of trainable parameters of encoder and decoder should be close.
s = reset_tf_session()
encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=32)
encoder.summary()
decoder.summary()
```

Layer (type)	Output Shape	Param #
=====		

input_1 (InputLayer)	(None, 32, 32, 3)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_3 (Conv2D)	(None, 8, 8, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 128)	0
conv2d_4 (Conv2D)	(None, 4, 4, 256)	295168
max_pooling2d_4 (MaxPooling2D)	(None, 2, 2, 256)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 32)	32800
=====		
Total params: 421,216		
Trainable params: 421,216		
Non-trainable params: 0		
=====		
Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	(None, 32)	0
dense_2 (Dense)	(None, 1024)	33792
reshape_1 (Reshape)	(None, 2, 2, 256)	0
conv2d_transpose_1 (Conv2DTr	(None, 4, 4, 128)	295040
conv2d_transpose_2 (Conv2DTr	(None, 8, 8, 64)	73792
conv2d_transpose_3 (Conv2DTr	(None, 16, 16, 32)	18464
conv2d_transpose_4 (Conv2DTr	(None, 32, 32, 3)	867
=====		
Total params: 421,955		
Trainable params: 421,955		
Non-trainable params: 0		
=====		

Convolutional autoencoder training. This will take **1 hour**. You're aiming at ~0.0056 validation MSE and ~0.0054 training MSE.

```
s = reset_tf_session()

encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=32)

inp = L.Input(IMG_SHAPE)
code = encoder(inp)
```

```

reconstruction = decoder(code)

autoencoder = keras.models.Model(inputs=inp, outputs=reconstruction)
autoencoder.compile(optimizer="adamax", loss='mse')

# we will save model checkpoints here to continue training in case of kernel death
model_filename = 'autoencoder.{0:03d}.hdf5'
last_finished_epoch = None

#### uncomment below to continue training from model checkpoint
#### fill `last_finished_epoch` with your latest finished epoch
# from keras.models import load_model
# s = reset_tf_session()
# last_finished_epoch = 4
# autoencoder = load_model(model_filename.format(last_finished_epoch))
# encoder = autoencoder.layers[1]
# decoder = autoencoder.layers[2]

autoencoder.fit(x=X_train, y=X_train, epochs=25,
                validation_data=[X_test, X_test],
                verbose=True,
                initial_epoch=last_finished_epoch or 0) #remove callbacks

```

Train on 11828 samples, validate on 1315 samples

```

Epoch 1/25
11828/11828 [=====] - 12s 991us/step - loss: 0.0126 -
Epoch 2/25
11828/11828 [=====] - 8s 643us/step - loss: 0.0078 -
Epoch 3/25
11828/11828 [=====] - 8s 639us/step - loss: 0.0071 -
Epoch 4/25
11828/11828 [=====] - 8s 646us/step - loss: 0.0070 -
Epoch 5/25
11828/11828 [=====] - 8s 649us/step - loss: 0.0068 -
Epoch 6/25
11828/11828 [=====] - 8s 642us/step - loss: 0.0067 -
Epoch 7/25
11828/11828 [=====] - 8s 648us/step - loss: 0.0066 -
Epoch 8/25
11828/11828 [=====] - 8s 643us/step - loss: 0.0065 -
Epoch 9/25
11828/11828 [=====] - 8s 641us/step - loss: 0.0064 -
Epoch 10/25
11828/11828 [=====] - 8s 642us/step - loss: 0.0062 -
Epoch 11/25
11828/11828 [=====] - 8s 678us/step - loss: 0.0061 -
Epoch 12/25
11828/11828 [=====] - 8s 661us/step - loss: 0.0060 -
Epoch 13/25
11828/11828 [=====] - 8s 645us/step - loss: 0.0059 -
Epoch 14/25
11828/11828 [=====] - 8s 641us/step - loss: 0.0058 -
Epoch 15/25
11828/11828 [=====] - 8s 651us/step - loss: 0.0057 -
Epoch 16/25
11828/11828 [=====] - 8s 642us/step - loss: 0.0056 -
Epoch 17/25

```

```

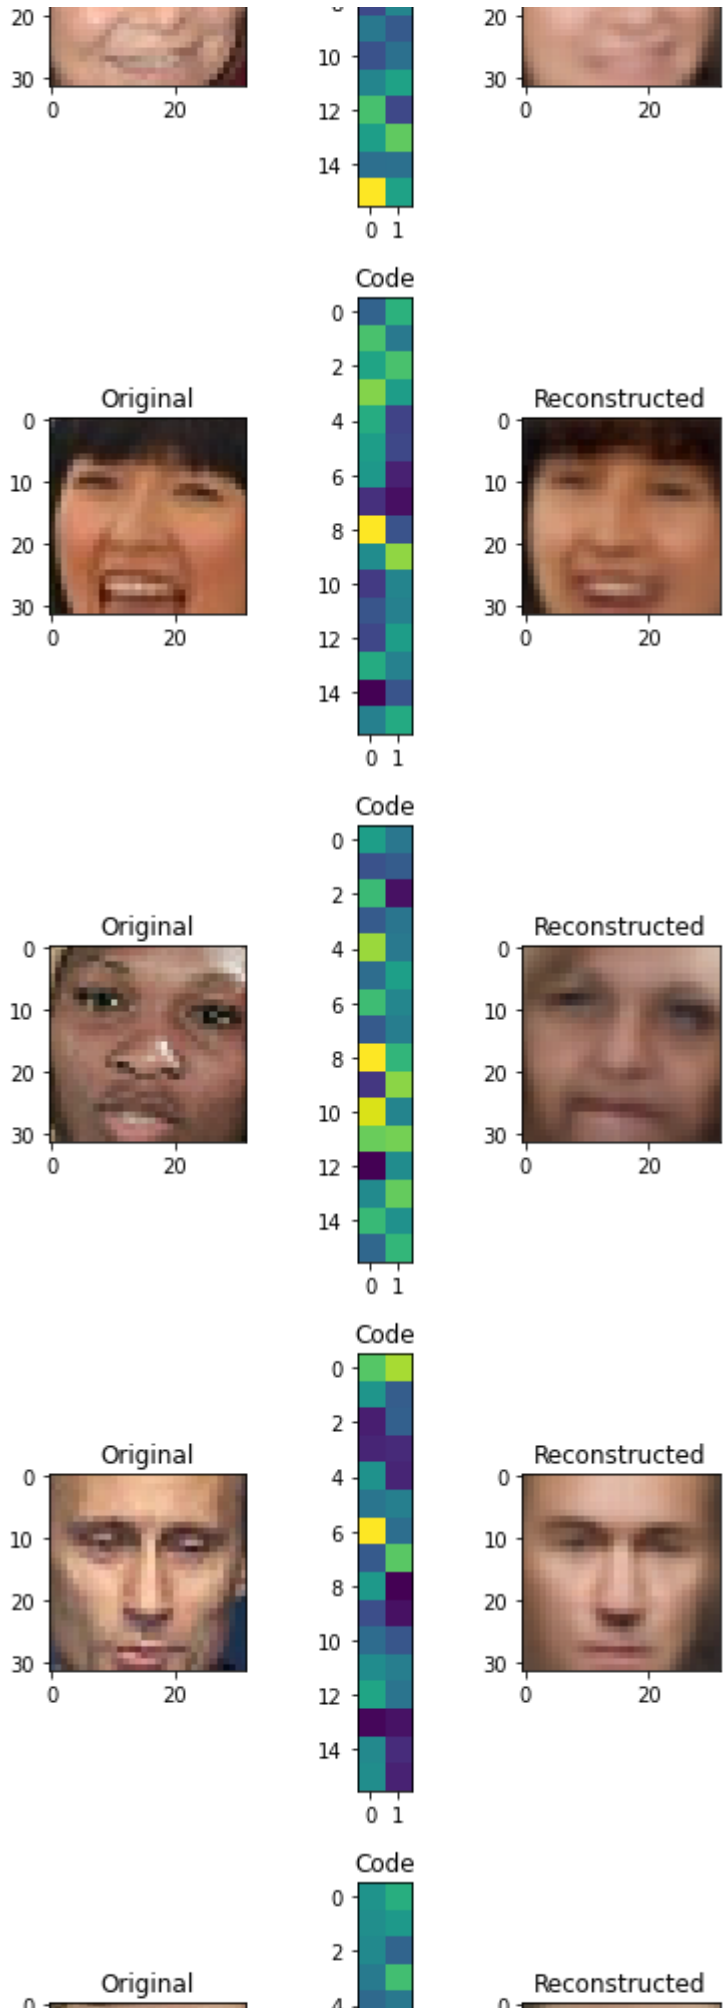
11828/11828 [=====] - 8s 645us/step - loss: 0.0056 -
Epoch 18/25
11828/11828 [=====] - 8s 644us/step - loss: 0.0055 -
Epoch 19/25
11828/11828 [=====] - 8s 648us/step - loss: 0.0055 -
Epoch 20/25
11828/11828 [=====] - 8s 646us/step - loss: 0.0054 -
Epoch 21/25
11828/11828 [=====] - 8s 645us/step - loss: 0.0054 -
Epoch 22/25
11828/11828 [=====] - 8s 642us/step - loss: 0.0053 -
Epoch 23/25
11828/11828 [=====] - 8s 646us/step - loss: 0.0053 -
Epoch 24/25
11828/11828 [=====] - 8s 651us/step - loss: 0.0052 -
Epoch 25/25
11828/11828 [=====] - 8s 650us/step - loss: 0.0052 -
<keras.callbacks.History at 0x7f5c20974cd0>

```

```

reconstruction_mse = autoencoder.evaluate(X_test, X_test, verbose=0)
print("Convolutional autoencoder MSE:", reconstruction_mse)
for i in range(5):
    img = X_test[i]
    visualize(img, encoder, decoder)

```


```
# save trained weights
encoder.save_weights("encoder.h5")
decoder.save_weights("decoder.h5")
```

```
# restore trained weights
s = reset_tf_session()

encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=32)
encoder.load_weights("encoder.h5")
decoder.load_weights("decoder.h5")

inp = L.Input(IMG_SHAPE)
code = encoder(inp)
reconstruction = decoder(code)

autoencoder = keras.models.Model(inputs=inp, outputs=reconstruction)
autoencoder.compile(optimizer="adamax", loss='mse')

print(autoencoder.evaluate(X_test, X_test, verbose=0))
print(reconstruction_mse)
```

```
0.005468181393219038
0.005468181393219038
```

▼ Submit to Coursera

```
from submit import submit_autoencoder
submission = build_deep_autoencoder(IMG_SHAPE, code_size=71)

# token expires every 30 min
COURSERA_TOKEN = 'sxBHLJ7pc2IWX55e' ### YOUR TOKEN HERE
COURSERA_EMAIL = 'e0321294@u.nus.edu' ### YOUR EMAIL HERE

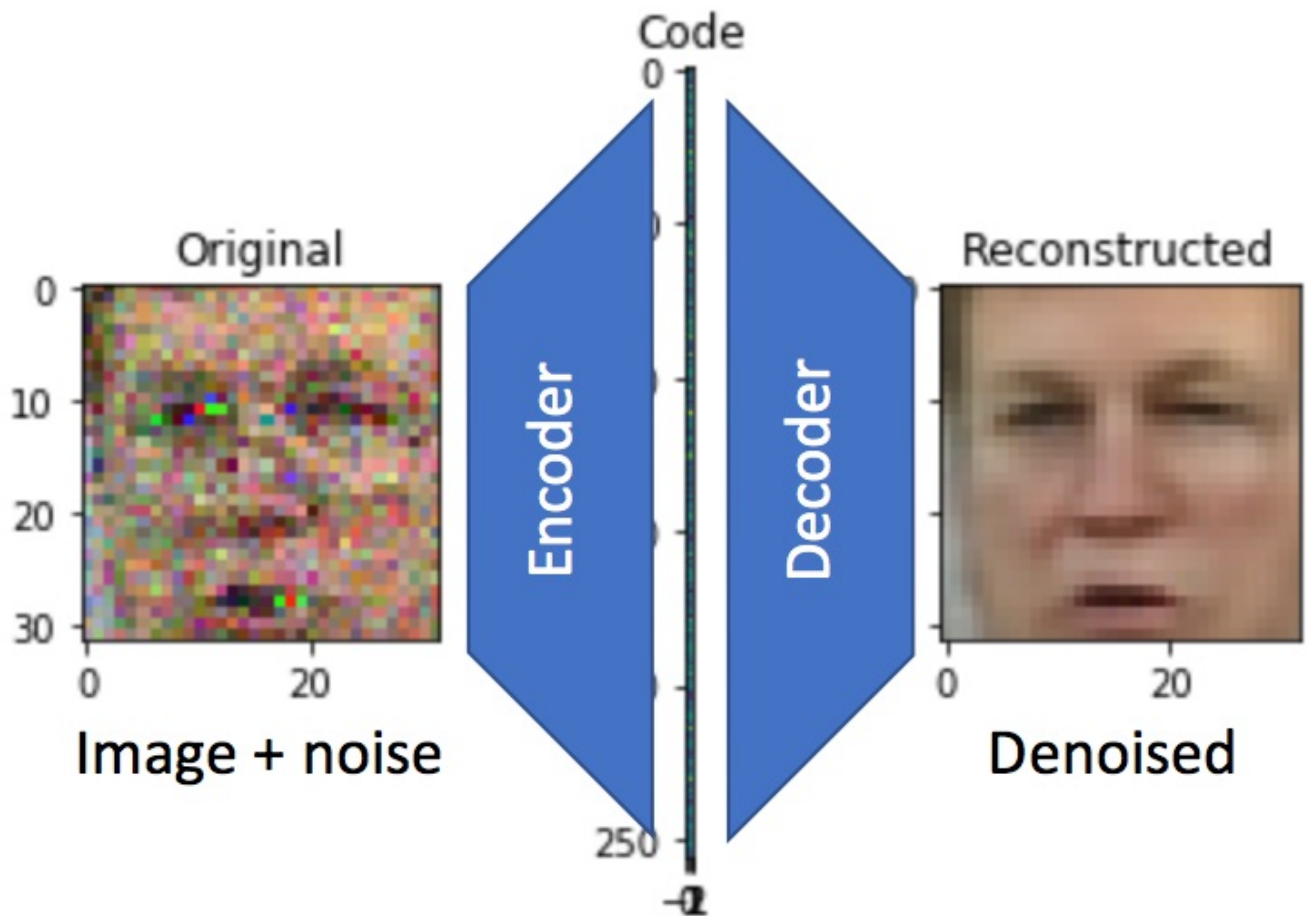
submit_autoencoder(submission, reconstruction_mse, COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

▼ Optional: Denoising Autoencoder

This part is **optional**, it shows you one useful application of autoencoders: denoising. You can run this code and make sure denoising works :)

Let's now turn our model into a denoising autoencoder:




We'll keep the model architecture, but change the way it is trained. In particular, we'll corrupt its input data randomly with noise before each epoch.

```
def apply_gaussian_noise(X, sigma=0.1):
    """
    adds noise from standard normal distribution with standard deviation sigma
    :param X: image tensor of shape [batch,height,width,3]
    Returns X + noise.
    """
    batch, height, width, ch = X.shape
    noise = np.random.normal(0, sigma, (height, width, ch)) ### YOUR CODE HERE ###
    return X + noise
```

```
# noise tests
theoretical_std = (X_train[:100].std()**2 + 0.5**2)**.5
our_std = apply_gaussian_noise(X_train[:100], sigma=0.5).std()
assert abs(theoretical_std - our_std) < 0.01, "Standard deviation does not match it
assert abs(apply_gaussian_noise(X_train[:100], sigma=0.5).mean() - X_train[:100].mea
```

```
# test different noise scales
plt.subplot(1,4,1)
show_image(X_train[0])
plt.subplot(1,4,2)
show_image(apply_gaussian_noise(X_train[:1], sigma=0.01)[0])
plt.subplot(1,4,3)
show_image(apply_gaussian_noise(X_train[:1], sigma=0.1)[0])
```

```
show_image(apply_gaussian_noise(X_train[:1],sigma=0.1)[0])
plt.subplot(1,4,4)
show_image(apply_gaussian_noise(X_train[:1],sigma=0.5)[0])
```



Training will take **1 hour**.

```
s = reset_tf_session()

# we use bigger code size here for better quality
encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=512)
assert encoder.output_shape[1:]==(512,), "encoder must output a code of required si

inp = L.Input(IMG_SHAPE)
code = encoder(inp)
reconstruction = decoder(code)

autoencoder = keras.models.Model(inp, reconstruction)
autoencoder.compile('adamax', 'mse')

for i in range(25):
    print("Epoch %i/25, Generating corrupted samples..."%(i+1))
    X_train_noise = apply_gaussian_noise(X_train)
    X_test_noise = apply_gaussian_noise(X_test)

    # we continue to train our model with new noise-augmented data
    autoencoder.fit(x=X_train_noise, y=X_train, epochs=1,
                    validation_data=[X_test_noise, X_test],
                    verbose=True)

    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
    11828/11828 [=====] - 8s 684us/step - loss: 0.0027 -
    Epoch 12/25, Generating corrupted samples...
    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
    11828/11828 [=====] - 8s 689us/step - loss: 0.0026 -
    Epoch 13/25, Generating corrupted samples...
    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
    11828/11828 [=====] - 8s 668us/step - loss: 0.0026 -
    Epoch 14/25, Generating corrupted samples...
    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
    11828/11828 [=====] - 8s 668us/step - loss: 0.0025 -
    Epoch 15/25, Generating corrupted samples...
    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
    11828/11828 [=====] - 8s 671us/step - loss: 0.0024 -
    Epoch 16/25, Generating corrupted samples...
    Train on 11828 samples, validate on 1315 samples
    Epoch 1/1
```

```

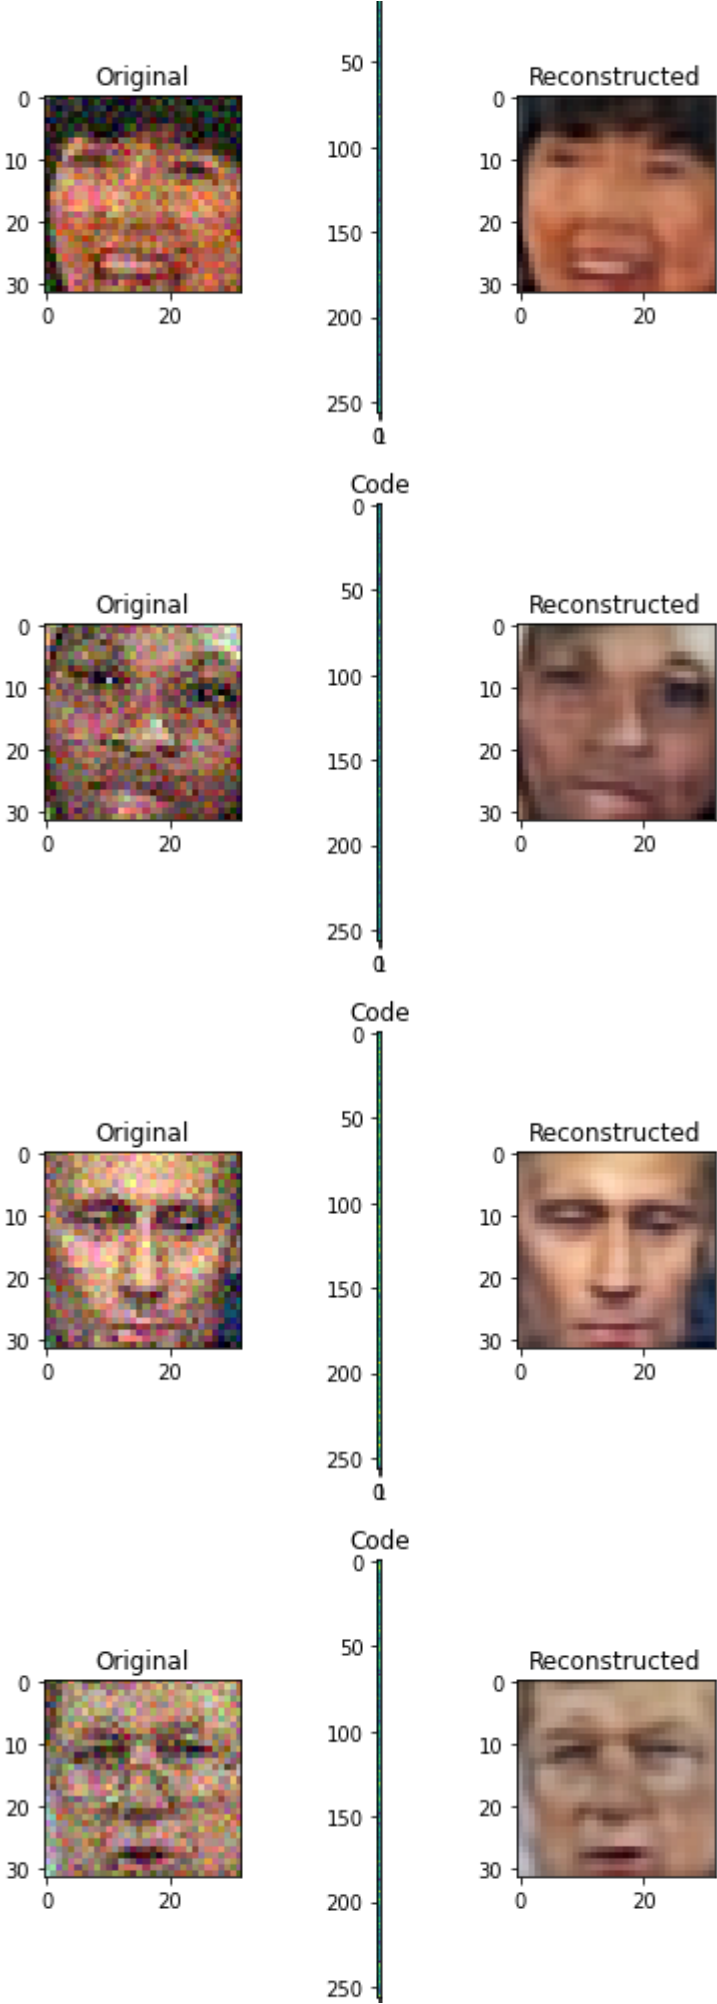
Epoch 1/1
11828/11828 [=====] - 8s 672us/step - loss: 0.0024 -
Epoch 17/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 672us/step - loss: 0.0023 -
Epoch 18/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 681us/step - loss: 0.0022 -
Epoch 19/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 680us/step - loss: 0.0022 -
Epoch 20/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 679us/step - loss: 0.0022 -
Epoch 21/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 680us/step - loss: 0.0021 -
Epoch 22/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 682us/step - loss: 0.0021 -
Epoch 23/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 696us/step - loss: 0.0020 -
Epoch 24/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 690us/step - loss: 0.0020 -
Epoch 25/25, Generating corrupted samples...
Train on 11828 samples, validate on 1315 samples
Epoch 1/1
11828/11828 [=====] - 8s 673us/step - loss: 0.0020 -

```

```

X_test_noise = apply_gaussian_noise(X_test)
denoising_mse = autoencoder.evaluate(X_test_noise, X_test, verbose=0)
print("Denoising MSE:", denoising_mse)
for i in range(5):
    img = X_test_noise[i]
    visualize(img, encoder, decoder)

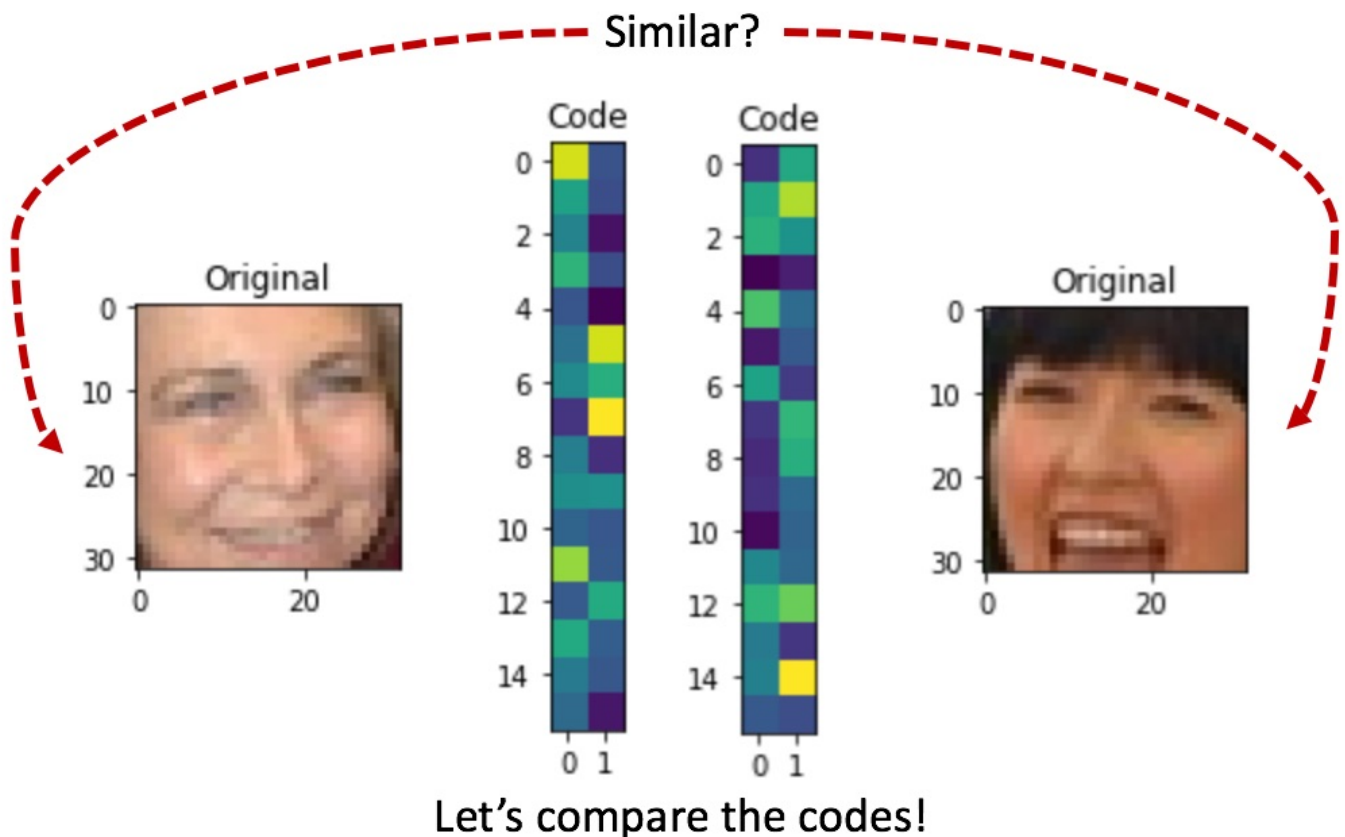
```



Optional: Image retrieval with autoencoders

So we've just trained a network that converts image into itself imperfectly. This task is not that useful in and of itself, but it has a number of awesome side-effects. Let's see them in action.

First thing we can do is image retrieval aka image search. We will give it an image and find similar images in latent space:



To speed up retrieval process, one should use Locality Sensitive Hashing on top of encoded vectors. This [technique](#) can narrow down the potential nearest neighbours of our image in latent space (encoder code). We will calculate nearest neighbours in brute force way for simplicity.

```
# restore trained encoder weights
s = reset_tf_session()
encoder, decoder = build_deep_autoencoder(IMG_SHAPE, code_size=32)
encoder.load_weights("encoder.h5")
```

```
images = X_train
codes = encoder.predict(images) ### YOUR CODE HERE: encode all images ###
assert len(codes) == len(images)
```



```
from sklearn.neighbors import NearestNeighbors
nei_clf = NearestNeighbors(metric="euclidean")
nei_clf.fit(codes)
```

```
NearestNeighbors(metric='euclidean')
```

```
def get_similar(image, n_neighbors=5):
    assert image.ndim==3,"image must be [batch,height,width,3]"

    code = encoder.predict(image[None])

    (distances,),(idx,) = nei_clf.kneighbors(code,n_neighbors=n_neighbors)

    return distances,idx
```

```
def show_similar(image):


    distances,neighbors = get_similar(image,n_neighbors=3)

    plt.figure(figsize=[8,7])
    plt.subplot(1,4,1)
    show_image(image)
    plt.title("Original image")

    for i in range(3):
        plt.subplot(1,4,i+2)
        show_image(neighbors[i])
        plt.title("Dist=%.3f"%distances[i])
    plt.show()
```

Cherry-picked examples:

```
# smiles
show_similar(X_test[247])
```



```
# ethnicity
show_similar(X_test[56])
```